

Parallel and Distributed Computing

CS3006 (BCS-6C/6D)

Lecture 12

Instructor: Dr. Syed Mohammad Irteza

Assistant Professor, Department of Computer Science, FAST

07 March, 2023

Previous Lecture

- OpenMP Locks
- Clauses
 - Private
 - Firstprivate
 - Lastprivate
 - If condition
- Scheduling – guided()
- Environment Variables:
 - OMP_NUM_THREADS
 - OMP_DYNAMIC
 - OMP_SCHEDULE
 - OMP_NESTED
- Computing Pi with Monte Carlo method

```
#pragma omp parallel shared(niter) private(i, x, y, z, chunk_size, seed) reduction(+:count) {

    num_threads = omp_get_num_threads();
    chunk_size = niter / num_threads;
    seed = omp_get_thread_num();
    #pragma omp master
    { printf("chunk_size=%ld\n", chunk_size); }
```

```
count=0;
for (i=0; i < chunk_size; i++) {
    //get random points
    x = (double) rand_r(&seed) / (double) RAND_MAX;
    y = (double) rand_r(&seed) / (double) RAND_MAX;
    z = ((x-0.5)*(x-0.5))+((y-0.5)*(y-0.5));
    //check to see if point is in unit circle
    if (z<0.25) {
        ++count;
    }
}
pi = ((double) count / (double) niter) * 4.0;
```

```
total number of allocated cores are:16
chunk_size=6250000
parallel Pi: 3.141515
Parallel time: 0.9560 seconds
Seq_Pi: 3.141745
Sequential time: 13.3521 seconds
speedup: 13.9669
```

Total points= 10 millions

```
total number of allocated cores are:16
chunk_size=62500000
parallel Pi: 3.141598
Parallel time: 8.5668 seconds
Seq_Pi: 3.141576
Sequential time: 132.0383 seconds
speedup: 15.4128
```

Total points= 100 millions

Computing Pi using the Monte Carlo method

(Parallel construct [parallel_pi.c])

More Detailed Discussion

- Full Example Online:
http://www.umsl.edu/~siegelj/cs4790/openmp/pimonti_omp.c.HTML
- Further Reading (optional):
 - <https://1drv.ms/p/s!Apc0G8okxWJ12jlUANAQsYO-JVdx?e=VixgYX> (just slide 1-9)
 - https://passlab.github.io/CSCE569/notes/lecture04-07_OpenMP.pdf
 - <https://www3.nd.edu/~zxu2/acms60212-40212/Lec-12-OpenMP.pdf>

Classic computing

- **Advantages:**

- Ease of programming (shared memory)
- Robustness
- Good chance of optimization

- **Disadvantages:**

- data protection from illegal operations
- low performance
 - single CPU, sequential computation
 - mitigated with multi-CPU systems and concurrent programming (thread, processes)
- requires physical access to the system (for usage)
 - mitigated with modem / network connections

Distributed System

- A distributed System is:


A collection of independent computers that appears to its users as a single coherent system

Distributed Systems

- Networks of computers are everywhere!
 - Mobile phone networks
 - Campus networks
 - In-car networks
 - On board networks in planes and trains
- “A system in which hardware or software components located at *networked* computers communicate and coordinate their actions only by *message passing*.” [Coulouris]
- “A distributed system is a collection of *independent* computers *that appear* to the users of the system as a single computer.” [Tanenbaum]

Reasons for Distributed Systems

- Functional Separation:
 - Existence of computers with different capabilities and purposes:
 - Clients and Servers
 - Data collection and data processing
- Inherent distribution:
 - Information:
 - Different information is created and maintained by different people (e.g., Web pages)
 - People
 - Computer supported collaborative work (virtual teams, engineering, virtual surgery)
 - Retail store and inventory systems for supermarket chains (e.g., Coles, Woolworths)



Reasons for Distributed Systems


- Power imbalance and load variation:
 - Distribute computational load among different computers.
- Reliability:
 - Long term preservation and data backup (replication) at different locations.
- Economies:
 - Sharing a printer by many users and reduce the cost of ownership.
 - Building a supercomputer out of a network of computers.



Distributed systems

- “A distributed system is a system designed to *support* the development of applications and services which can *exploit* a physical architecture consisting of multiple *autonomous* processing elements that *do not share primary memory* but cooperate by sending *asynchronous messages* over a communications network”
 - [Blair and Stefani, 1998]

Distributed Systems

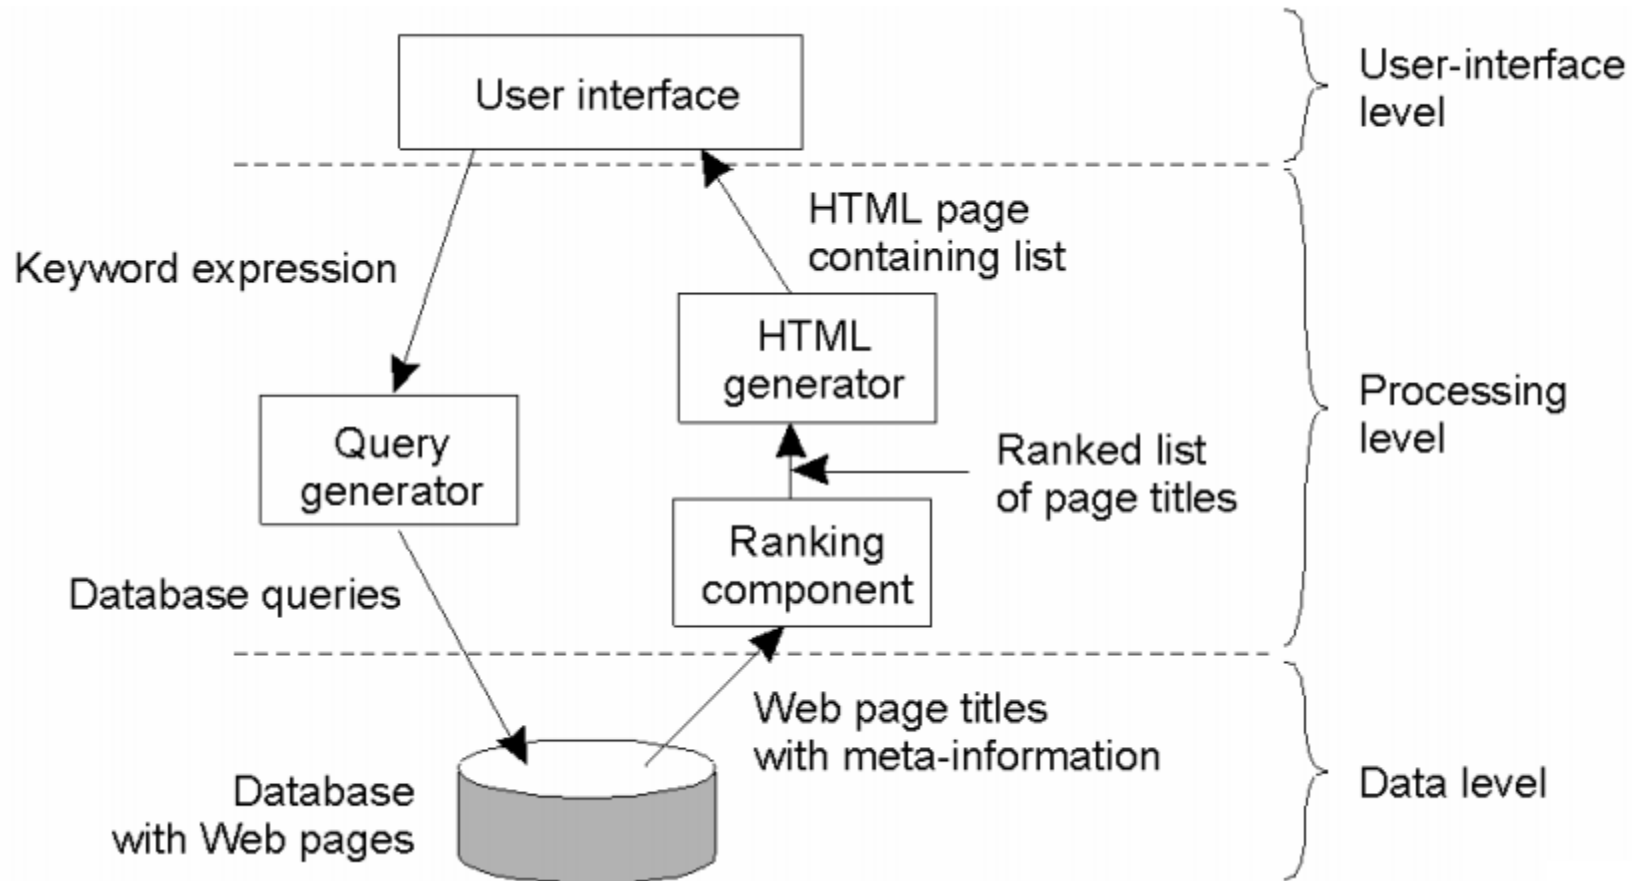


- Each system has its local data (private)
- Multiple address spaces
- Concurrent computation over different CPUs
- multiple computation flows

Examples of distributed systems

- Internet: network of networks
 - Global access to everyone
 - Huge size
 - No single authority
- Intranets
 - Protected access
 - Single authority
- E-learning
- Online gaming
- Mobile ad-hoc networks

Search engine



- <http://www.ece.rutgers.edu/~parashar>

Advantages of Distributed systems

- High performance
 - Several CPUs
- Good scalability
 - Increasing the number of CPUs is easier than increasing the performance of a single CPU
- Data protected from illegal operations
 - Disjoint memory spaces, accessible only by their respective programs
- Network access
 - user physical presence not required

Trends in distributed systems

- **Pervasive network technology**

- modern Internet is a vast interconnected collection of computer networks of many different types
- range of types increasing all the time
- devices can be connected at any time and at any place

- **Mobile and ubiquitous computing**

- integration of small and portable computing devices into distributed systems
- Laptops, PDAs, smart watches, smart embedded devices in appliances
- location-aware or context-aware computing

Trends in distributed systems

- presence of computers everywhere only becomes useful when they can communicate with one another
- it is more desirable for users to control their washing machine or their entertainment system from their phone or a 'universal remote control' device in the home
- **Spontaneous interoperation**
 - associations between devices are routinely created and destroyed
 - Challenges: should be fast and convenient

Trends in distributed systems

- Distributed computing as a utility
 - resources are provided by service suppliers and effectively rented rather than owned by the end user
 - Cloud Computing
- Focus on resource sharing
 - we routinely share hardware resources such as printers, files, and resources with more specific functionality such as search engines
 - Client/Server computing
 - Peer to Peer computing

Forms of Distributed Computing

Cluster Computing

“A computer cluster is a group of loosely coupled computers that work together closely so that in many respects it can be viewed as though it were a single computer. Clusters are commonly connected through fast local area networks.”

Forms of Distributed Computing

Grid Computing

- “Grid is an infrastructure that involves the integrated and collaborative use of Computers, networks, databases and scientific instruments owned and managed by multiple organizations”
- Grid Computing is Biased towards the solution of Scientific problems

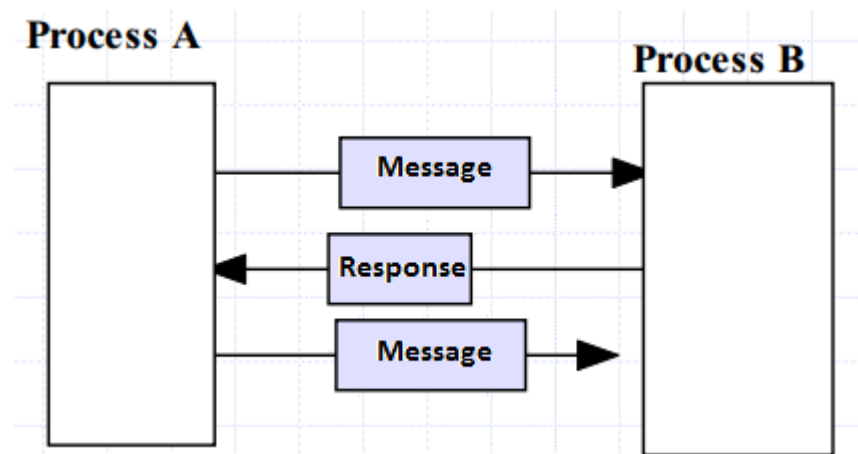
Forms of Distributed Computing

Cloud Computing

A cloud is defined as a set of Internet-based application, storage and computing services sufficient to support most users' needs, thus enabling them to largely or totally dispense with local data storage and application software

Message passing paradigm

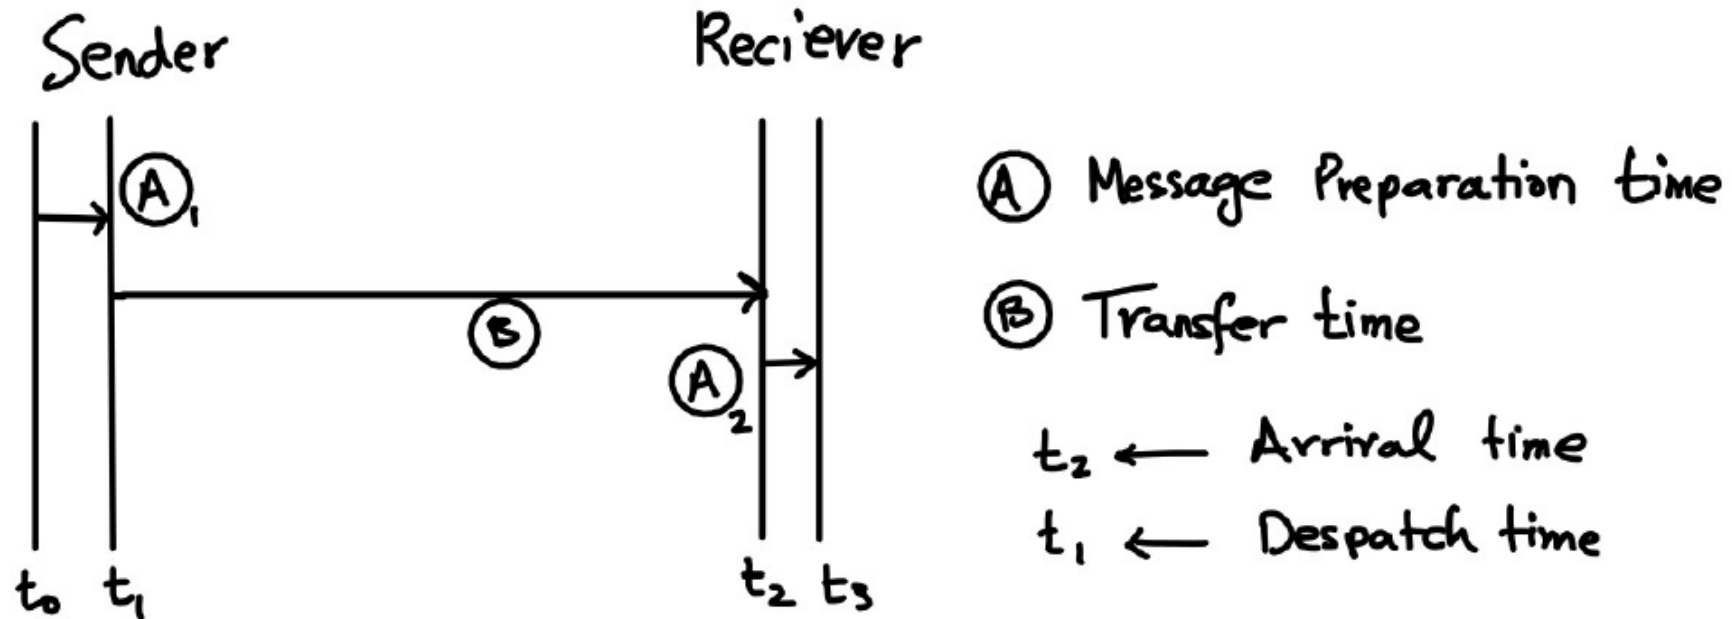
- Most fundamental paradigm for distributed applications



- A *process* sends a message representing a *request*
- The message is delivered to a *receiver*, which processes the *request*, and sends a message in *response*
- In turn, the *reply* may trigger a further *request*, which leads to a subsequent *reply*, and so forth

Communication Model

- Message preparation & Message transfer time



- Arrival latency = transfer time = $t_2 - t_1$

Parameters for measuring arrival latency

Bandwidth : Data Volume (r) sent / unit time

e.g 400 mbits / second \Rightarrow 1 mbit transfer = $\frac{1s}{400}$


$$1 \text{ unit transfer} = 1 / r$$

- $T_w = 1 \text{ unit transfer time} = 1/r$
- $T_s = \text{preparation time} = t_1 - t_0$
- $T_h = \text{Arrival latency} = t_2 - t_1$
- $T_r = \text{Receiver handling time} = t_3 - t_2$
- Size of message = m
- No. of hops/nodes, the message is to be sent across = n

- Total time for sending a message of size m across n hops:


$$t = t_s + (t_w m + t_h + t_r) n$$

- Example:



Example: $t = t_s + (t_w m + t_h + t_r) n$

- Bandwidth = 2 Mbps (megabits per second)
- m (message size) = 100 megabits
- $n = 2$



Example: $t = t_s + (t_w m + t_h + t_r) n$

- Bandwidth = 2 Mbps (megabits per second)
- m (message size) = 100 megabits
- $n = 2$


$T_w = 1$ unit transfer time = $\frac{1}{2}$ (in terms of Mbits)

$$T_s = t_1 - t_0 = 2 - 0 = 2$$

$$T_h = t_2 - t_1 = 5 - 2 = 3$$

$$T_r = t_3 - t_2 = 6 - 5 = 1$$

$$\begin{aligned} t &= t_s + (t_w m + t_h + t_r) n \\ &= 2 + (\frac{1}{2} * 100 + 3 + 1) * 2 = 110s \end{aligned}$$




Architectures

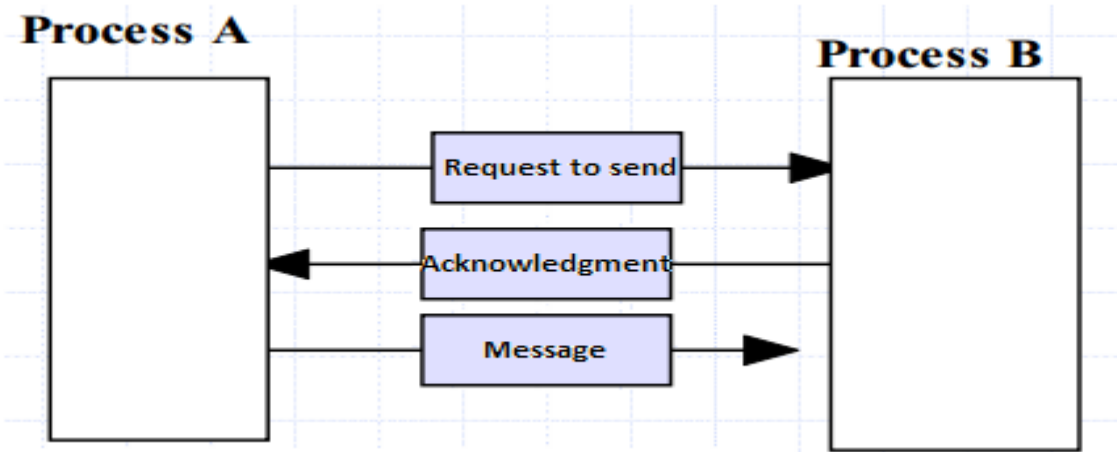
- Client-Server (C/S) architecture
 - Asymmetric architecture
 - Server position is determined *a priori*
- Peer-to-Peer (P2P) architecture
 - Symmetric architecture
 - Every node can play both the client and server roles (simultaneously or at different times)

Client/Server paradigm

- Most *widely used paradigm* for network applications
- *Server* plays the role of a *service provider* which waits passively for the arrival of request
- *Client* issues specific *requests* to the server and waits for its response



Synchronous message passing

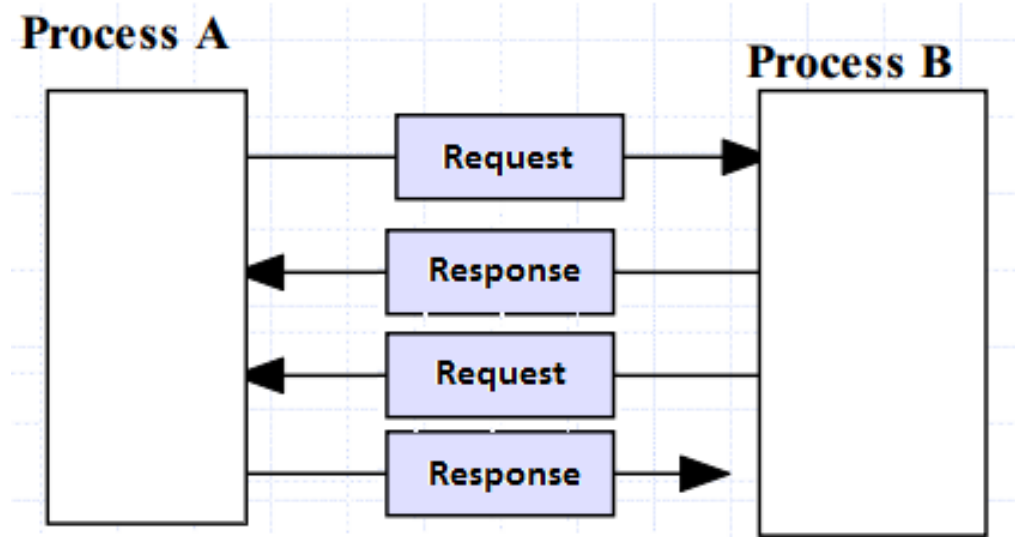


- Routines *return* when message transfer *completed*
- Send routine *waits* until complete message can be accepted by the receiving process before *sending* the message
- Receive routine *waits* until the message it is expecting *arrives*

Asynchronous Message Passing

- Routines *do not wait* for actions to complete before returning
- Local *storage/message queue* is required to hold messages
- Advantage: *Faster*, since it allows processes to move forward
- Disadvantage: *Less reliable* and must be used with care
- Message buffer needed between source and destination to hold messages
- Send routine *wait* in case the buffer space is exhausted

Peer to Peer system architecture



- In the peer-to-peer paradigm, the participating processes play equal roles, with *equivalent capabilities and responsibilities* (hence the term “peer”)
- *Each* participant may issue a *request* to another participant and receive a *response*

Peer to Peer system architecture

- Computer resources and services are *directly exchanged* between computer systems
- Resources and services include the *exchange of information, processing cycles, cache storage, and disk storage for files*
- Computers that have traditionally been used solely as clients communicate directly among themselves and *can act as both clients and servers*, assuming whatever role is most efficient for the network
- The peer-to-peer paradigm is more appropriate for applications such as *instant messaging, peer-to-peer file transfers, video conferencing, and collaborative work*

P2P systems: some history

- [1.3 P2P Systems - Rui Pan's Blog \(gitbook.io\)](#)
 - Napster
 - Gnutella
 - BitTorrent

Distributed Systems: Issues

- *Programming Complexity*
 - How the various programs communicate together?
 - Which data format on the various network nodes?
 - Need to define (application) protocols
 - Operations synchronization may lead to delay and slowing down
- *Scarce Robustness*
 - Higher chance of errors/faults
- *Hard to Optimize*
 - Lack of a global view

Distributed systems

- “A distributed system is a system designed to *support* the development of applications and services which can *exploit* a physical architecture consisting of multiple *autonomous* processing elements that *do not share primary memory* but cooperate by sending *asynchronous messages* over a communications network”
 - [Blair and Stefani, 1998]

Operating System

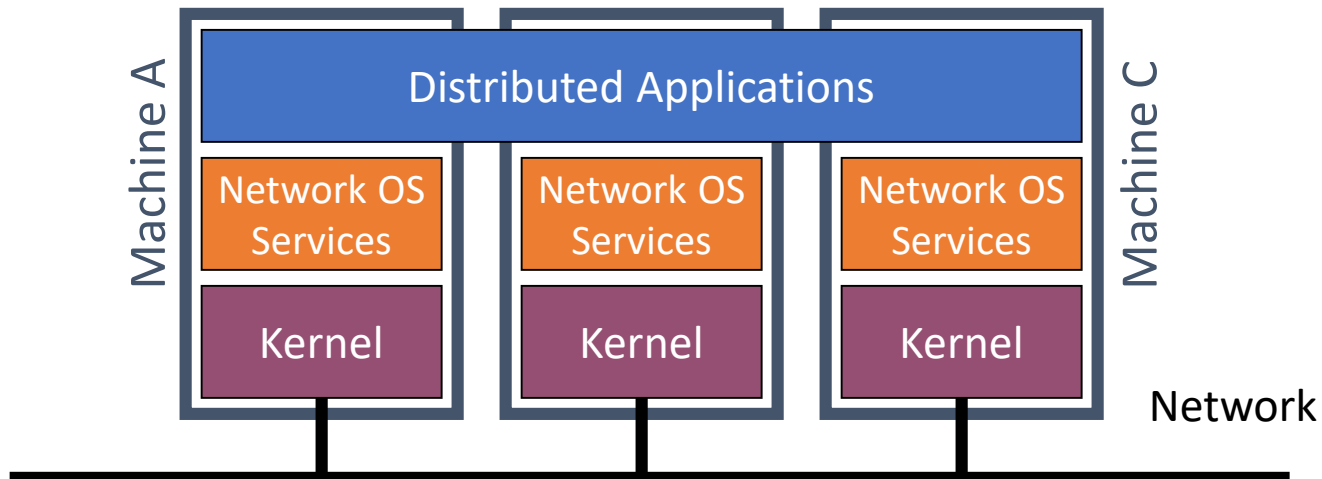
- A program that manages the computer hardware
- Provides a basis for application programs and acts as an intermediary between the user and the hardware
- Offers a reasonable way to solve the problem of creating a usable computing system

Types of Operating Systems:

- Linux and Unix
- Microsoft Windows
- Mac OS X

Network Operating Systems

- Distribution **not hidden** by OS
- But a number of **services** offered to support distribution
 - remote login (ssh), remote copying (scp)
 - distributed file systems (samba, nfs)
- Well adapted to **heterogeneous DS**



Features of Network OS

- Provides basic operating system features; *support for processors, protocols, automatic hardware detection*, support *multi-processing of applications*
- Security features; *authentication, authorization, access control*
- Provides files, print, web services, back-up and *replication services*
- Supports Internetworking such as *routing* and *WAN ports*
- User management and support for *login/logoff, remote access, system management*

Network and Operating System Security

- OS: *system must protect itself*
- List of possible breaches is almost endless
- Runaway process could constitute an accidental denial-of-service attack
- Travels over private leased lines, shared lines like the internet, wireless connections, or dial-up lines
- Intercepting these data could be harmful as breaking into a computer
- Interruption of communications could constitute a remote denial-of-service attack

Distributed OSs

- **Completely hide** the distribution: “single system image”
- Each individual node holds a specific software subset of the global aggregate operating system.
- The user has the **illusion** to use one single multiprocessor machine (symmetric multiprocessing, SMP)
- Essentially used for **homogeneous** cluster of machines interconnected with **high performance networks**

